ISABELLE QIAN

Berkeley, California | (408) 442-4009 | isabelle.qian@gmail.com | linkedin.com/in/isabelle-athena-qian

EDUCATION

University of California, Berkeley Graduation: Spring 2027

B.A. Computer Science (Honors) | *B.A.* Philosophy | *B.A.* German Studies | *Minor:* Public Policy *GPA:* 3.979/4.0

Relevant Coursework: Engineering Responsible Systems, Agentic AI, Machine Learning, Principles & Techniques of Data Science, Probability Theory, User Interface Design, Internet Architecture, Language after LLMs

AI GOVERNANCE AND POLICY RESEARCH

UC Berkeley School of Information | *Research Assistant to Professor Deirdre Mulligan* June 2025 – Present

- Compiled 100+ federal, state, and international AI governance policies for Eric Hysen, former DHS Chief AI Officer.
- Built a database of California laws (2024–2026) on public-sector AI use; mapped 2025 AI bill sponsors.
- Built a tracker of 14 lawsuits and enforcement actions against frontier AI developers, coding cited safety evidence.
- Co-developed UPLBench, a benchmark built from court decisions that audits LLM legal advice across five high-stakes domains; built the annotation interface, annotator training, and adjudication workflows.
- Sole-authored a July 2026 oversight memo on external red-teaming of GPT-4o: found public records for 87 of 97 testers, and none of five psychology-tagged testers had published on self-harm or harm to minors.
- Researched 30+ U.S. data infrastructures for a cross-agency field guide under Denice Ross, former U.S. Chief Data Scientist.

UC Berkeley School of Law | *Research Assistant to Professor Colleen Chien* September 2025 – May 2026

- Researched intellectual property, AI, and algorithmic pricing policy, plus criminal justice reform through the Paper Prisons Initiative; designed research methods and coordinated a team of research assistants.

International Association for Safe and Ethical AI (IASEAI) | *Associate Project Manager* January 2026 – Present

- Coordinated 12+ workshops for ~400 participants at IASEAI'26 at UNESCO House, Paris (1,500+ attendees), running participant communications and a live registration and waitlist tracker.
- Drafted the IASEAI Code of Conduct and the conference playbook, translating governance principles into community rules.

PUBLICATIONS AND PRESENTATIONS

- **A Decision-Theoretic Approach for Managing Misalignment**, IASEAI'26, February 2026 (co-author; accepted and presented). Framework for deciding when to delegate to AI agents, weighing value alignment against epistemic accuracy.
- **UPLBench: Auditing LLM Legal Advice**, AI for Law Scholars Conference 2026 (co-author; accepted and presented).

TECHNICAL AND QUANTITATIVE PROJECTS

Real-time AI Legislation Tracker Python, LegiScan, SerpApi, Gemini, Streamlit

- Built an ETL pipeline and cached Streamlit dashboard tracking legislative text and media coverage for 50+ U.S. state and federal AI bills, with Gemini prompts that turn statutes into compliance summaries.
- Extracted risk frameworks and stakeholder positions from bill text to cut manual legislative desk research.

White Agent: WebSearchRAG and AIPolicyBench Python | Fall 2025, CS194 Agentic AI

- Built a retrieval agent for government AI policy documents and AIPolicyBench, a 300-item LLM-as-a-judge evaluation scoring correctness against hallucination; reached ~2.7x baseline accuracy (20.3% vs. 7.6%).

PRODUCT, ADVISING, AND TEACHING

GenAI Learning Project, Brilliance of Berkeley January – May 2026

- Designed multi-modal GenAI exercises for a 13-student pilot in a 744-student course to test responsible use of LLM outputs.
- Interviewed participants across disciplines to iterate on each exercise; co-authored the public course report.

Student Advisor | *UC Berkeley Office of Undergraduate Education (VPUE)* September 2024 – Present

- Co-authored the undergraduate AI guidance distributed to all incoming students; shaped faculty-student GenAI dialogues.

Student Representative | *Student Integrated Systems Executive Committee* May 2025 – Present

- Gather student feedback for monthly decisions on digital infrastructure, service integration, and governance practices.

Undergraduate Student Instructor | *UC Berkeley Mathematics* July 2024 – August 2026

- Taught 13 Math 1 sections of ~25 students each across three terms, covering 13 pre-calculus topics; adapted materials and Zoom instruction for accessibility across varied academic backgrounds.

HONORS

- **Dean's List** (2023–2026); **Samuel Silver Memorial Scholarship** (EECS, 2026); **Vukasin Leadership Award** (Residential Life, 2024–25); **Best First Year Mentor and Best Cohort** (CDSS, 2025).

SKILLS

- *Focus areas:* Responsible AI, AI governance and policy, AI trust and safety evaluation, privacy and data governance, algorithmic accountability, policy research, user-centered design, accessibility
- *Technical:* Python, R, SQL, Java, C; NumPy, Pandas, scikit-learn, Hugging Face, Streamlit, Git; RAG, LLM-as-a-judge evals
- *Languages:* English and Mandarin Chinese (native); German (professional working proficiency)